


CAN201: Introduction to Networking

Lecture 7 - The Network Layer: Routing Algorithm



Lecturers: Dr. Gordon Boateng & Dr. Fei Cheng

Important Information

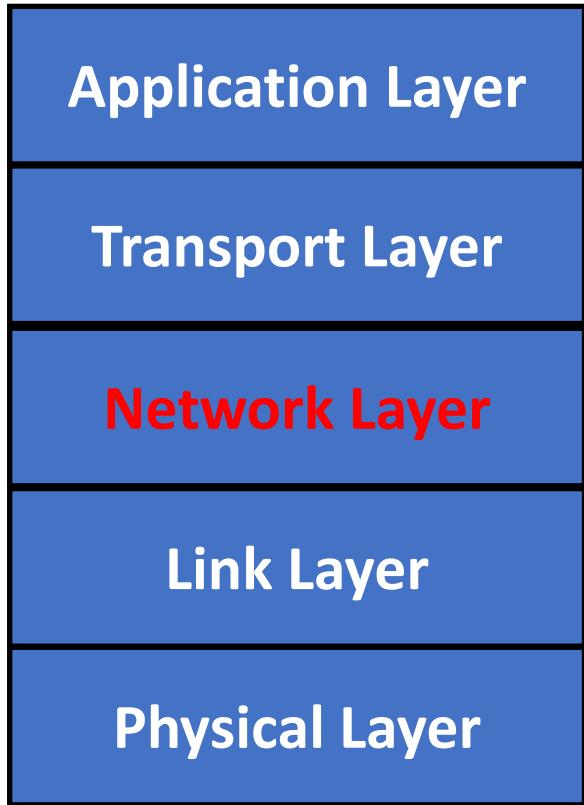
■ Contact:

- Email: Gordon.Boateng@xjtu.edu.cn
- Office No.: SC554A

■ Office Hours (Strictly via appointment)

- Tuesday: 14:00-15:00
- Wednesday: 14:00-15:00

Recap: Top-Down Approach




Network Layer: Data Plane

- Overview of Network Layer
- The Internet Protocol: IPv4, Addressing, NAT
IPv6
- Generalized Forwarding and SDN
- Routing Protocols

Network-layer Services and Protocols

- Transport segment from sending host to receiving host
 - **sender**: encapsulates segments into datagrams, passes to link layer
 - **receiver**: delivers segments to transport layer protocol
- network layer protocols in *every Internet device*: hosts, routers
- **router**:
 - examines header fields in all IP datagrams passing through it
 - moves datagrams from input ports to output ports to transfer datagrams along end-end path



Two Key Network-layer Functions

network-layer functions:

- *forwarding*: move packets from a router's input link to appropriate router output link
- *routing*: determine route taken by packets from source to destination
 - *routing algorithms*

analogy: taking a trip

- *forwarding*: process of getting through single interchange
- *routing*: process of planning trip from source to destination



forwarding




routing

Network layer: Data plane, Control plane

Data plane:

- *local*, per-router function
- determines how datagram arriving on router input port is forwarded to router output port



Control plane

- *network-wide* logic
- determines how datagram is routed among routers along end-end path from source host to destination host
- two control-plane approaches:
 - *traditional routing algorithms*: implemented in routers
 - *software-defined networking (SDN)*: implemented in (remote) servers


Q1: What is the difference between control plane and data plane?

Network Layer: Data Plane

- Overview of Network Layer
- The Internet Protocol:
 - IPv4, addressing (part 1)
 - NAT, IPv6 (part 2)
- Generalized Forwarding and SDN
- Routing Protocols

Network Layer: Internet


host, router network layer functions:



IP Addressing: Introduction

- **Interface:** connection between host/router and physical link
 - routers typically have multiple interfaces
 - host typically has one or two interfaces (e.g., wired Ethernet, wireless 802.11)
- **IP address:** 32-bit identifier associated with each host or router *interface*

Q2: if a router has 5 interfaces, how many IP addresses does it have?



dotted-decimal IP address notation:

223.1.1.1 =


11011111	00000001	00000001	00000001
223	1	1	1

IP Addressing: Introduction (cont'd)

Q: How are interfaces actually connected?

A: wired
Ethernet interfaces
connected by
Ethernet switches

Q3: What is the difference
between a router and a switch?



A: wireless WiFi interfaces
connected by WiFi base station


Subnets

- *What is a subnet ?*

- device interfaces that can physically reach each other **without passing through an intervening router**

- IP addresses have two parts:

- **subnet part:** devices in same subnet have common high order bits
- **host part:** remaining low order bits




network consisting of 3 subnets

Subnets (cont'd)

Recipe for defining subnets:

- detach each interface from its host or router, creating “islands” of isolated networks
- each isolated network is called a *subnet*




*Subnet 3:
223.1.3.0/24*

subnet mask: /24
(high-order 24 bits: subnet part of IP address)

Subnets (cont'd)

- where are the subnets?
- what are the /24 subnet addresses?



IP addressing: Network Classes

Classful addressing:

- The network portion of an IP address were constrained to be 8, 16, or 24 bits in length.
- Subnets with 8-, 16-, and 24-bit subnet addresses were known as class A, B and C networks.
- It became problematic
 - A class B (/16) subnet supporting $2^{16} - 2 = 65,534$ hosts, which is too large

Under classful addressing, an organization with 10,000 hosts, which subnet class should be allocated? What is the problem?

This led to a rapid depletion of the class B address space and poor utilization of the assigned address space.

IP addressing: CIDR

CIDR: Classless InterDomain Routing (pronounced “cider”)

- subnet portion of address can have arbitrary length
- address format: **a.b.c.d/x**, where x is # bits in subnet portion of address



IP addresses: how to get one?

Two questions:

1. Q: How does a *host* get IP address within its network (host part of address)?
2. Q: How does a *network* get IP address for itself (network part of address)

How does *host* get IP address?

- hard-coded by sysadmin in config file (e.g., Windows: control-panel->network->configuration->tcp/ip->properties, /etc/rc.config in UNIX)
- **DHCP: Dynamic Host Configuration Protocol:** dynamically get address from a server

DHCP: Dynamic Host Configuration Protocol

Goal: Allow host to *dynamically* obtain its IP address from network server when it “joins” network


- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/on)
- support for mobile users who join/leave network

DHCP overview:

- host broadcasts **DHCP discover** msg [optional]
- DHCP server responds with **DHCP offer** msg [optional]
- host requests IP address: **DHCP request** msg
- DHCP server sends address: **DHCP ack** msg

DHCP Client-Server Scenario

DHCP server: 223.1.2.5




DHCP: more than IP addresses


DHCP can return more than just allocated IP address on subnet:

- address of first-hop router for client
- name and IP address of DNS sever
- network mask (indicating network versus host portion of address)

DHCP Client-Server Scenario (cont'd)



DHCP: Example



- Connecting laptop will use DHCP to get IP address, address of first-hop router, address of DNS server.
- DHCP REQUEST message encapsulated in UDP, encapsulated in IP, encapsulated in Ethernet
- Ethernet frame broadcast (dest: FFFFFFFFFFFF) on LAN, received at router running DHCP server
- Ethernet demux'ed to IP demux'ed, UDP demux'ed to DHCP

IP addresses: how to get one?

Q: how does *network* get subnet part of IP address?

A: gets allocated portion of its provider ISP's address space

ISP's block	<u>11001000</u> <u>00010111</u> <u>00010000</u> <u>00000000</u> 200.23.16.0/20
-------------	--

ISP can then allocate out its address space in, say, 8 blocks:

Organization 0	<u>11001000</u> <u>00010111</u> <u>00010000</u> <u>00000000</u> 200.23.16.0/23
----------------	--

Organization 1	<u>11001000</u> <u>00010111</u> <u>00010010</u> <u>00000000</u> 200.23.18.0/23
----------------	--

Organization 2	<u>11001000</u> <u>00010111</u> <u>00010100</u> <u>00000000</u> 200.23.20.0/23
----------------	--

...

.....

.....


.....

Organization 7	<u>11001000</u> <u>00010111</u> <u>00011110</u> <u>00000000</u> 200.23.30.0/23
----------------	--

Hierarchical addressing: route aggregation


Hierarchical addressing allows efficient advertisement of routing information:

The ability to use a single prefix to advertise multiple networks is often referred to as **address aggregation** or **route aggregation**.




Hierarchical addressing: more specific routes

- Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



Hierarchical addressing: more specific routes

- Organization 1 moves from Fly-By-Night-ISP to ISPs-R-Us
- ISPs-R-Us now advertises a more specific route to Organization 1



IP addressing: last words ...

Q: How does an ISP get block of addresses?

A: ICANN: Internet Corporation for Assigned Names and Numbers

<http://www.icann.org/>

- allocates IP addresses, through 5 regional registries (RRs) (who may then allocate to local registries)
- manages DNS root zone, including delegation of individual TLD (.com, .edu , ...) management

Q: Are there enough 32-bit IP addresses?


- ICANN allocated last chunk of IPv4 addresses to RRs in 2011
- NAT (next) helps IPv4 address space exhaustion
- IPv6 has 128-bit address space

Network Layer: Data Plane

- Overview of Network Layer
- The Internet Protocol:
 - IPv4, addressing (part 1)
 - NAT, IPv6 (part 2)
- Generalized Forwarding and SDN
- Routing Protocols

NAT: Network Address Translation

NAT: All devices in the local network share just **one** IPv4 address as far as the outside world is concerned



all datagrams *leaving* local network have *same* source NAT IP address: 138.76.29.7, but *different* source port numbers

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

NAT: Network Address Translation (cont'd)

- all devices in local network have 32-bit addresses in a “private” IP address space (10/8, 172.16/12, 192.168/16 prefixes) that can only be used in local network
- advantages:
 - just **one** IP address needed from provider ISP for ***all*** devices
 - can change addresses of host in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - security: devices inside local network not directly addressable, visible by outside world


NAT: Network Address Translation (cont'd)

implementation: NAT router must (transparently):

- outgoing datagrams: replace (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
 - remote clients/servers will respond using (NAT IP address, new port #) as destination address
- remember (in NAT translation table) every (source IP address, port #) to (NAT IP address, new port #) translation pair
- incoming datagrams: replace (NAT IP address, new port #) in destination fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: Network Address Translation (cont'd)

2: NAT router changes datagram source address from 10.0.0.1, 3345 to 138.76.29.7, 5001, updates table



NAT: Network Address Translation (cont'd)

- NAT has been controversial:
 - routers “should” only process up to layer 3
 - address “shortage” should be solved by IPv6
 - violates end-to-end argument (port # manipulation by network-layer device)
 - NAT traversal: what if client wants to connect to server behind NAT?
- but NAT is here to stay:
 - extensively used in home and institutional nets, 4G/5G cellular nets

IPv6: Motivation

- **initial motivation:** 32-bit IPv4 address space would be completely allocated
- Additional motivation:
 - Header format helps speed processing/forwarding
 - Header changes to facilitate QoS


IPv6 datagram format:

- Fixed-length 40-byte header
- No fragmentation allowed (at intermediate routers, see “packet too big” new ICMP message type)

IPv6 Datagram Format

priority: identify priority among datagrams in flow

128-bit IPv6 addresses




flow label: identify datagrams in same "flow." (concept of "flow" not well defined).

What's missing (compared with IPv4):

- no checksum (to speed processing at routers)
- no fragmentation/reassembly
- no options (available as upper-layer, next-header protocol at router)


Transition from IPv4 to IPv6

- Not all routers can be upgraded simultaneously
 - no “flag days”
 - how will network operate with mixed IPv4 and IPv6 routers?
- **Tunneling:** IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers (“packet within a packet”)
 - tunneling used extensively in other contexts (4G/5G)




Tunneling and Encapsulation

Ethernet connecting
two IPv6 routers:




IPv4 network
connecting two
IPv6 routers




Tunneling and Encapsulation

Ethernet connecting
two IPv6 routers:



IPv4 tunnel
connecting two
IPv6 routers




Tunneling

logical view:



physical view:




Network Layer: Data Plane

- Overview of Network Layer
- The Internet Protocol: IPv4, Addressing, NAT
IPv6
- **Generalized Forwarding and SDN**
- Routing Protocols


Flow Table Abstraction

- **flow:** defined by header field values (in link-, network-, transport-layer fields)
- **generalized forwarding:** simple packet-handling rules
 - **match:** pattern values in packet header fields
 - **actions:** for matched packet: drop, forward, modify, matched packet or send matched packet to controller
 - **priority:** disambiguate overlapping patterns
 - **counters:** #bytes and #packets




Flow Table Abstraction

- **flow:** defined by header fields
- **generalized forwarding:** simple packet-handling rules
 - Pattern: match values in packet header fields
 - Actions: for matched packet: drop, forward, modify, matched packet or send matched packet to controller
 - Priority: disambiguate overlapping patterns
 - Counters: #bytes and #packets



OpenFlow: Flow Table Entries



OpenFlow: Examples

Destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	51.6.0.8	*	*	*	*	port6

IP datagrams destined to IP address 51.6.0.8 should be forwarded to router output port 6

Firewall:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	*	*	*	*	*	22	drop

Block (do not forward) all datagrams destined to TCP port 22 (ssh port #)

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
*	*	*	*	*	*	128.119.1.1	*	*	*	*	*	drop

Block (do not forward) all datagrams sent by host 128.119.1.1

OpenFlow: Examples

Layer 2 destination-based forwarding:

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	VLAN Pri	IP Src	IP Dst	IP Prot	IP ToS	TCP s-port	TCP d-port	Action
-------------	---------	---------	----------	---------	----------	--------	--------	---------	--------	------------	------------	--------

* * 22:A7:23:
11:E1:02 * * * * * * * * * * * port3

layer 2 frames with destination MAC address 22:A7:23:11:E1:02 should be forwarded to output port 3

OpenFlow Abstraction

- **match+action:** abstraction unifies different kinds of devices

Router

- *match:* longest destination IP prefix
- *action:* forward out a link

Switch

- *match:* destination MAC address
- *action:* forward or flood

Firewall

- *match:* IP addresses and TCP/UDP port numbers
- *action:* permit or deny

NAT

- *match:* IP address and port
- *action:* rewrite address and port

Network Layer: Data Plane

- Overview of Network Layer
- The Internet Protocol: IPv4, Addressing, NAT
IPv6
- Generalized Forwarding and SDN
- Routing Protocols

Routing Protocols


Routing protocol goal: determine “good” paths (equivalently, routes), from sending host to receiving host, through network of routers

- **Path:** sequence of routers packets will traverse in going from given initial source host to given final destination host
- “good”: least “cost”, “fastest”, “least congested”
- **Routing:** a “top-10” networking challenge!

Graph Abstraction of the Network

A graph is often used to formulate routing problems..

graph: $G = (N, E)$




N = set of nodes (routers) = { u, v, w, x, y, z }

E = set of edges (links) = { (u,v), (u,x), (v,x), (v,w), (v,y), (w,x), (w,y), (w,z), (y,z) }

aside: graph abstraction is useful in other network contexts, e.g., P2P, where N is set of peers and E is set of TCP connections

Graph Abstraction: Link Costs




$c(x,x')$ = cost of link (x,x')
e.g., $c(w,z) = 5$

cost could always be 1, or
inversely related to
bandwidth,
or directly related to
congestion.

cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

key question: what is the least-cost path between u and z ?
routing algorithm: algorithm that finds that least cost path

Routing Algorithm Classification



A Link-State Routing Algorithm - Dijkstra

Dijkstra's algorithm

- Net topology and link costs known to all nodes
 - Accomplished via “link state broadcast”
 - All nodes have same info
- Computes least cost paths from one node (“source”) to all other nodes
 - Gives *forwarding table* for that node
- Iterative: after k iterations, know least cost path to k dest.’s

Notations:

- $c(x,y)$: link cost from node x to y ; $= \infty$ if not direct neighbors
- $D(v)$: current value of cost of path from source to dest. v
- $p(v)$: predecessor node along path from source to dest. v
- N' : set of nodes whose least cost path is definitively known

Dijkstra's Algorithm

```
1 Initialization:
2    $N' = \{u\}$ 
3   for all nodes  $v$ 
4     if  $v$  adjacent to  $u$ 
5       then  $D(v) = c(u,v)$ 
6     else  $D(v) = \infty$ 
7
8 Loop
9   find  $w$  not in  $N'$  such that  $D(w)$  is a minimum
10  add  $w$  to  $N'$ 
11  update  $D(v)$  for all  $v$  adjacent to  $w$  and not in  $N'$  :
12     $D(v) = \min(D(v), D(w) + c(w,v))$ 
13  /* new cost to  $v$  is either old cost to  $v$  or known
14    shortest path cost to  $w$  plus cost from  $w$  to  $v$  */
15 until all nodes in  $N'$ 
```

Notation:

- $c(x,y)$: link cost from node x to y ; $= \infty$ if not direct neighbors
- $D(v)$: current value of cost of path from source to dest. v
- $p(v)$: predecessor node along path from source to v
- N' : set of nodes whose least cost path definitively known


Dijkstra's Algorithm: Example 1

computing **routing table** in u:

Step	N'	D(v)	D(w)	D(x)	D(y)	D(z)
		p(v)	p(w)	p(x)	p(y)	p(z)
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uwx	6,w			11,w	14,x
3	uwxv			10,v	14,x	
4	uwxvy				12,y	
5	uwxvzy					


notes:

- ❖ construct shortest path tree by tracing predecessor nodes
- ❖ ties can exist (can be broken arbitrarily)



Dijkstra's Algorithm: Example 2 - result

resulting shortest-path tree from u:



resulting **forwarding table** in u:

destination	link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

Distance Vector Algorithm

- Bellman-Ford equation (dynamic programming)

let

$d_x(y) :=$ cost of least-cost path from x to y

then


$$d_x(y) = \min_v \{ c(x, v) + d_v(y) \}$$

cost to neighbor v

cost from neighbor v to destination y

\min_v is taken over all neighbors v of x

Bellman-Ford Example



B-F equation:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

From u's perspective estimated:

clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

Node achieving minimum is the next hop in shortest path, used in forwarding table

Distance Vector Algorithm

Key idea:

- From time-to-time, each node sends its own distance vector estimate to neighbors
- When x receives new DV estimate from neighbor, it updates its own DV using B-F equation:

$$D_x(y) \leftarrow \min_v \{c(x,v) + D_v(y)\} \text{ for each node } y \in N$$

- ❖ Under minor, natural conditions, the estimate $D_x(y)$ converge to the actual least cost $d_x(y)$

Distance Vector Algorithm

Iterative, asynchronous:


each local iteration caused by:

- Local link cost change
- DV update message from neighbor

Distributed:

- Each node notifies neighbors *only* when its DV changes
 - Its neighbors then further notify their neighbors if necessary

Each node:



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞


	x	y	z
x	0	2	3
y	2	0	1
z	7	1	0

node y table

	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

node z table

	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0



time

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

from	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

node y table

from	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

node z table

from	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0

time

