

CAN201: Introduction to Networking

Lecture 9 - The Link Layer 1



Lecturers: Dr. Gordon Boateng and Dr. Fei Cheng

Important Information

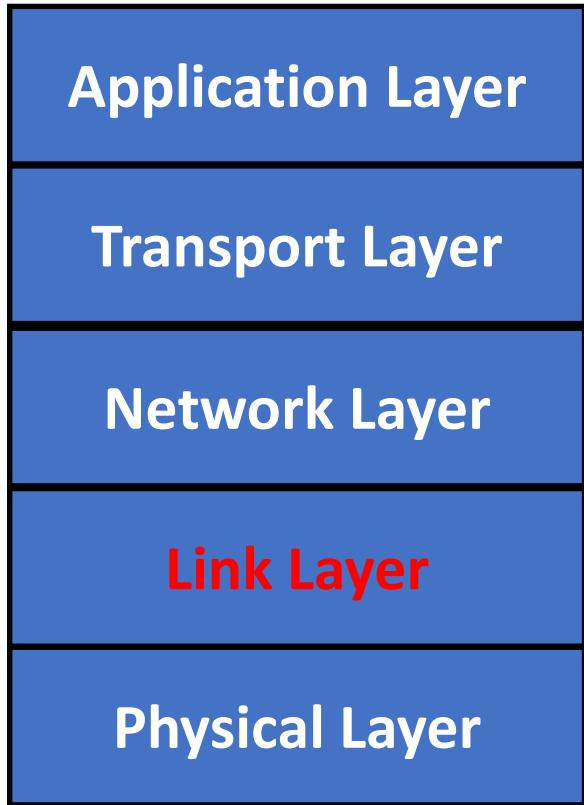
■ Contact:

- Email: Gordon.Boateng@xjtu.edu.cn
- Office No.: SC554A

■ Office Hours (Strictly via appointment)

- Tuesday: 14:00-15:00
- Wednesday: 14:00-15:00

Recap: Top-Down Approach



The Link Layer


- Introduction to the Link Layer
 - Error-Detection and -Correction Techniques
 - Multiple Access Links and Protocols

Link layer: Introduction

terminology:

- hosts, routers: nodes
- communication channels that directly connect physically adjacent nodes: links
 - wired , wireless
- layer-2 packet: *frame*, encapsulates datagram

link layer has responsibility of transferring datagram from one node to physically adjacent node over a link



Data Unit

Application Layer

Transport Layer

Network Layer

Link Layer

Physical Layer

?

Datagram

Frame

Message

Segment

Bit (or Signal)

Link Layer: Context

- datagram transferred by different link protocols over different links:
 - e.g., WiFi on first link, Ethernet on next link
- each link protocol provides different services
 - e.g., may or may not provide reliable data transfer over link

transportation analogy:

- trip from Suzhou to Osaka
 - limo: Suzhou to Shanghai Airport
 - plane: Shanghai Airport to Tokyo
 - train: Tokyo to Osaka
- tourist = datagram
- transport segment = communication link
- transportation mode = link-layer protocol
- travel agent = routing algorithm


Link Layer: Services

- **framing, link access:**

- encapsulate datagram into frame, adding header, trailer
- channel access if shared medium
- “MAC” addresses in frame headers identify source, destination (different from IP address!)

- **reliable delivery between adjacent nodes**

- we already know how to do this!
- seldom used on low bit-error links
- wireless links: high error rates
 - Q: why both link-level and end-end reliability?



Link Layer: Services (cont'd)

- **flow control:**

- pacing between adjacent sending and receiving nodes

- **error detection:**


- errors caused by signal attenuation, noise.
 - receiver detects errors, signals retransmission, or drops frame

- **error correction:**

- receiver identifies *and corrects* bit error(s) without retransmission


- **half-duplex and full-duplex:**

- with half duplex, nodes at both ends of link can transmit OR receive, but not at same time
 - with full duplex, end and receive simultaneously




Host Link-Layer Implementation

- In each and every host
- Link layer implemented in “Adapter” or on a chip
 - Ethernet card, Wifi 802.11 card, or chipset
 - For link and physical layers
- Attached into host’s system (motherboard) buses
 - USB / PCI / Thunderbolt ...
- Hardware / software / firmware



Adaptors Communicating



sending side:

- encapsulates datagram in frame
- adds error checking bits, reliable data transfer, flow control, etc.

receiving side:

- looks for errors, reliable data transfer, flow control, etc.
- extracts datagram, passes to upper layer at receiving side


The Link Layer

- Introduction to the Link Layer
- **Error-detection and -correction Techniques**
- Multiple Access Links and Protocols

Error Detection

EDC: error detection and correction bits (e.g., redundancy)

D: data protected by error checking, may include header fields



Error detection not 100% reliable!

- protocol may miss some errors, but rarely
- larger EDC field yields better detection and correction

Three techniques for detecting errors

- Parity Checks**

- basic ideas (simple, detect odd or even numbers of errors)

- Checksum**

- used in transport layer (but link layer can use variants)

- Cyclic Redundancy Checks (CRC)**

- used in link layer in an adapter (most powerful)

Parity Checking

single bit parity:

- detect single bit errors

0111000110101011	1
------------------	---

← d data bits → |
 parity bit

Even parity: set parity bit so there is an even number of 1's


Odd parity: set parity bit so there is an odd number of 1's

At receiver:

- compute parity of $d+1$ received bits, if not even, then error detected
- can detect odd number of bit flips (if two bits flip, it may fail)

E.g. 1: (Even parity) D=1000001; detect 1-bit error and 2-bit error

2-D parity:



- detect two-bit errors (adds row and column parity bits)
- detect *and correct* single bit errors without retransmission!



no errors:	1 0 1 0 1 1 (even)	detected and correctable single-bit error:
	1 1 1 1 0 0	
	0 1 1 1 0 1	
	<hr/> <hr/> <hr/>	
	1 0 1 0 1 0	
	(odd)	

E.g. 2: (odd parity) D=4x7 data block; check parity

1 0 1 0 1 1	parity error
1 0 1 1 0 0	
0 1 1 1 0 1	
1 0 1 0 1 0	
	↓ parity error

Internet Checksum

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

sender:

- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- **checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - not equal - error detected
 - equal - no error detected. *But maybe errors nonetheless?* More later

Internet checksum: 1s complement

Example: add two 16-bit integers

	1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
	1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
<hr/>																
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1
<hr/>																
sum	1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum	0	1	0	0	0	1	0	0	0	1	0	0	0	1	1	

Note: wraparound - when adding numbers, a carryout from the most significant bit needs to be added to the result

Cyclic redundancy check (CRC)

- more powerful error-detection coding
- widely used in practice (Ethernet, 802.11 WiFi, ATM)
- all CRC calculations are done in modulo-2 arithmetic (without carries in addition or borrows in subtraction).
- modulo-2 Arithmetic
 - Addition
 - Subtraction
 - Bitwise exclusive-or (XOR)


$$1011 + 0101 = ?$$

$$1011 - 0101 = ?$$

$$1011 \text{ XOR } 0101 = ?$$

Cyclic Redundancy Check (CRC)

- more powerful error-detection coding
- **D**: data bits (given, think of these as a binary number)
- **G**: bit pattern (generator), of $r+1$ bits (given, specified in CRC standard)



sender: Given D and $G(r+1)$ bits, append r bits R such that $D \cdot 2^r \text{ XOR } R$ is divisible by G ($\text{mod } 2$)

receiver: knows G , divides $\langle D, R \rangle$ by G . If remainder not equal to zero, error detected!

- can detect single-bit-error, double-bit error, all burst errors less than $r+1$ bits
- widely used in practice (Ethernet, 802.11 WiFi)

Cyclic Redundancy Check (CRC): Example

Sender wants to compute R such that:

$$D \cdot 2^r \text{ XOR } R = nG$$

... or equivalently (XOR R both sides):


$$D \cdot 2^r = nG \text{ XOR } R$$

... which says:

if we divide $D \cdot 2^r$ by G, we want remainder R to satisfy:

$$R = \text{remainder} \left[\frac{D \cdot 2^r}{G} \right]$$

algorithm for computing R



The Link Layer

- Introduction to the Link Layer
- Error-detection and -correction Techniques
- **Multiple Access Links and Protocols**

Multiple Access Links, Protocols

two types of “links”:

- point-to-point
 - point-to-point link between Ethernet switch, host
 - PPP for dial-up access
- broadcast (shared wire or medium)
 - old-school Ethernet
 - upstream HFC in cable-based access network
 - 802.11 wireless LAN, 4G/5G, satellite




shared wire (e.g.,
cabled Ethernet)



shared radio: 4G/5G



shared radio: WiFi



shared radio: satellite



humans at a cocktail party
(shared air, acoustical)

Multiple Access Protocols

- single shared broadcast channel
- two or more simultaneous transmissions by nodes: interference
 - *collision* if node receives two or more signals at the same time

multiple access protocol

- distributed algorithm that determines how nodes share channel, i.e., determine when node can transmit
- communication about channel sharing must use channel itself!
 - no out-of-band channel for coordination

An Ideal Multiple Access Protocol

given: multiple access channel (MAC) of rate R bps

Requirements:

1. when one node wants to transmit, it can send at rate R .
2. when M nodes want to transmit, each can send at average rate R/M
3. fully decentralized:
 - no special node to coordinate transmissions
 - no synchronization of clocks, slots
4. simple

MAC Protocols: Taxonomy

three broad classes:

- **channel partitioning**

- divide channel into smaller “pieces”
(time slots, frequency, code)
- allocate piece to node for exclusive
use

- **random access**

- channel not divided, allow collisions
- “recover” from collisions


- **“taking turns”**

- nodes take turns, but nodes with
more to send can take longer turns

Channel Partitioning MAC Protocols: TDMA

TDMA: time division multiple access


- access to channel in “rounds”
- each station gets fixed length slot (length = packet transmission time) in each round
- unused slots go idle
- example: 6-station LAN, 1,3,4 have packets to send, slots 2,5,6 idle



Channel Partitioning MAC Protocols: FDMA

FDMA: frequency division multiple access

- channel spectrum divided into frequency bands
- each station assigned fixed frequency band
- unused transmission time in frequency bands go idle
- example: 6-station LAN, 1,3,4 have packet to send, frequency bands 2,5,6 idle



Random Access Protocols

- when node has packet to send
 - transmit at full channel data rate R
 - no *a priori* coordination among nodes
- two or more transmitting nodes:
“collision”
- **random access protocol** specifies:
 - how to detect collisions
 - how to recover from collisions (e.g., via delayed retransmissions)
- examples of random access MAC protocols:
 - ALOHA, slotted ALOHA
 - CSMA, CSMA/CD, CSMA/CA

CSMA (Carrier Sense Multiple Access)

simple **CSMA**: listen before transmit:


- if channel sensed idle: transmit entire frame
- if channel sensed busy: defer transmission
- human analogy: don't interrupt others!

CSMA/CD: CSMA with *collision detection*

- collisions *detected* within short time
- colliding transmissions aborted, reducing channel wastage
- collision detection easy in wired, difficult with wireless
- human analogy: the polite conversationalist


CSMA: Collisions

- collisions can *still* occur with carrier sensing:
 - propagation delay means two nodes may not hear each other's just-started transmission
- **collision:** entire packet transmission time wasted
 - distance & propagation delay play role in determining collision probability



CSMA/CD:

- CSMA/CD reduces the amount of time wasted in collisions
 - transmission aborted on collision detection



“Taking Turns” MAC Protocols

channel partitioning MAC protocols:

- **Pro:** share channel *efficiently* and *fairly* at high load
- **Con:** inefficient at low load: delay in channel access, $1/N$ bandwidth allocated even if only 1 active node!

random access MAC protocols

- **Pro:** efficient at low load: single node can fully utilize channel
- **Con:** high load: collision overhead


“taking turns” protocols

- look for best of both worlds!

“Taking Turns” MAC Protocols (cont’d)

polling:


- centralized controller “invites” other nodes to transmit in turn
- typically used with “dumb” devices
- concerns:
 - polling overhead
 - latency
 - single point of failure (master)
 - Bluetooth uses polling



“Taking Turns” MAC Protocols (cont’d)

token passing:

- control *token* message explicitly passed from one node to next, sequentially
 - transmit while holding token
- concerns:
 - token overhead
 - latency
 - single point of failure (token)



Link layer, LANs: Roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
 - addressing, ARP
 - Ethernet
 - switches
 - VLANs


MAC Addresses

- 32-bit IP address:
 - *network-layer* address for interface
 - used for layer 3 (network layer) forwarding
 - e.g.: 128.119.40.136
 - MAC (or LAN or physical or Ethernet) address:
 - function: used “locally” to get frame from one interface to another physically-connected interface (same subnet, in IP-addressing sense)
 - 48-bit MAC address (for most LANs) burned in NIC ROM, also sometimes software settable
 - e.g.: 1A-2F-BB-76-09-AD
- hexadecimal (base 16) notation
(each “numeral” represents 4 bits)*

MAC Addresses

each interface on LAN

- has unique 48-bit **MAC** address
- has a locally unique 32-bit IP address (as we've seen)




MAC Addresses

- MAC address allocation administered by IEEE
- manufacturer buys portion of MAC address space (to assure uniqueness)
- analogy:
 - MAC address: like Social Security Number
 - IP address: like postal address
- MAC flat address: portability
 - can move interface from one LAN to another
 - recall IP address *not* portable: depends on IP subnet to which node is attached

ARP: Address Resolution Protocol

Question: how to determine interface's MAC address, knowing its IP address?



ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:
<IP address; MAC address; TTL>
- TTL (Time To Live): time after which address mapping will be forgotten (typically 20 min)


ARP Protocol in Action

example: A wants to send datagram to B

- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address

- 1 A broadcasts ARP query, containing B's IP addr
- destination MAC address = FF-FF-FF-FF-FF-FF
 - all nodes on LAN receive ARP query


ARP table in A		
IP addr	MAC addr	TTL



ARP protocol in Action (cont'd)

example: A wants to send datagram to B


- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



ARP protocol in Action (cont'd)

example: A wants to send datagram to B


- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



Routing to Another Subnet: Addressing

walkthrough: sending a datagram from *A* to *B* via *R*

- focus on addressing – at IP (datagram) and MAC layer (frame) levels
- assume that:
 - A knows B's IP address
 - A knows IP address of first hop router, R (how?)
 - A knows R's MAC address (how?)




Link layer, LANs: Roadmap

- introduction
- error detection, correction
- multiple access protocols
- **LANs**
 - addressing, ARP
 - **Ethernet**
 - switches
 - VLANs

Ethernet

“dominant” wired LAN technology:

- first widely used LAN technology
- simpler, cheap
- kept up with speed race: 10 Mbps – 400 Gbps
- single chip, multiple speeds (e.g., Broadcom BCM5761)




*Metcalfe's Ethernet
sketch*


Ethernet: Physical Topology

- **bus:** popular through mid 90s
 - all nodes in same collision domain (can collide with each other)
- **switched:** prevails today
 - active link-layer 2 *switch* in center
 - each “spoke” runs a (separate) Ethernet protocol (nodes do not collide with each other)

bus: coaxial cable



switched



Ethernet Frame Structure


sending interface encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**



preamble:

- used to synchronize receiver, sender clock rates
- 7 bytes of 10101010 followed by one byte of 10101011

Ethernet Frame Structure (cont'd)



- **addresses**: 6 byte source, destination MAC addresses
 - if adapter receives frame with matching destination address, or with broadcast address (e.g., ARP packet), it passes data in frame to network layer protocol
 - otherwise, adapter discards frame
- **type**: indicates higher layer protocol
 - mostly IP but others possible, e.g., Novell IPX, AppleTalk
- **CRC**: cyclic redundancy check at receiver
 - error detected: frame is dropped

Ethernet: Unreliable, Connectionless

- **connectionless:** no handshaking between sending and receiving NICs
- **unreliable:** receiving NIC does not send ACKs or NAKs to sending NIC
 - data in dropped frames recovered only if initial sender uses higher layer rdt (e.g., TCP), otherwise dropped data lost
- Ethernet's MAC protocol: unslotted **CSMA/CD with binary backoff**